

# Database Management Systems

Lecture 8

Evaluating Relational Operators

Query Optimization

- running example - schema
  - Students (SID: integer, SName: string, Age: integer)
  - Courses (CID: integer, CName: string, Description: string)
  - Exams (SID: integer, CID: integer, EDate: date, Grade: integer)
- Students
  - every record has 50 bytes
  - there are 80 records / page
  - 500 pages
- Courses
  - every record has 40 bytes
  - there are 100 records / page
  - 1 page
- Exams
  - every record has 40 bytes
  - there are 100 records / page
  - 1000 pages

## Projection

- $\Pi_{\text{SID, CID}}(\text{Exams})$

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- to implement projection:
  - eliminate:
    - unwanted columns
    - duplicates
- projection algorithms - *partitioning* technique:
  - sorting
  - hashing

## Projection Based on Sorting

- step 1
  - scan  $E \Rightarrow$  set of tuples containing only desired attributes ( $E'$ )
  - cost:
    - scan  $E$ :  $M$  I/Os
    - write temporary relation  $E'$ :  $T$  I/Os
      - $T$  depends on: number of columns and their sizes,  $T$  is  $O(M)$
- step 2
  - sort tuples in  $E'$
  - sort key: all columns
  - cost:  $O(T \log T)$  (also  $O(M \log M)$ )
- step 3
  - scan sorted  $E'$ , compare adjacent tuples, eliminate duplicates
  - cost:  $T$
- total cost:  $O(M \log M)$

## Projection Based on Sorting

### \* example

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- scan Exams: 1000 I/Os
- size of tuple in E': 10 bytes

=> cost of writing temporary relation E': 250 I/Os

- available buffer pages: 20
  - E' can be sorted in 2 passes
  - sorting cost:  $2 * 2 * 250 = 1000$  I/Os
- final scan of E' - cost: 250 I/Os

=> total cost:  $1000 + 250 + 1000 + 250 = 2500$  I/Os

\* E – record size = 40 bytes \*                      \* 1000 pages \*                      \* 100 records / page\*

## Projection Based on Sorting

### \* example

```
SELECT DISTINCT E.SID, E.CID  
FROM Exams E
```

- scan Exams: 1000 I/Os
- size of tuple in E': 10 bytes

=> cost of writing temporary relation E': 250 I/Os

- available buffer pages: 257
  - E' can be sorted in 1 pass
  - sorting cost:  $2 * 1 * 250 = 500$  I/Os
- final scan of E' - cost: 250 I/Os

=> total cost:  $1000 + 250 + 500 + 250 = 2000$  I/Os


\* E – record size = 40 bytes \*                      \* 1000 pages \*                      \* 100 records / page\*

## Projection Based on Sorting

- improvement
  - adapt the sorting algorithm to do projection with duplicate elimination
    - modify pass 0 of External Merge Sort: eliminate unwanted columns
      - read in B pages from E
      - write out  $(T/M) * B$  internally sorted pages of E'
        - refinement: write out  $2*B$  internally sorted pages of E' (on average)
      - tuples in runs - smaller than input tuples
    - modify merging passes: eliminate duplicates
      - number of result tuples is smaller than number of input tuples

\* minimize the number of runs - optional

- external merge sort
  - $N$  pages in the file,  $B$  buffer pages  $\Rightarrow \lceil N/B \rceil$  runs of  $B$  pages each
- improvement
  - algorithm known to produce runs of approximately  $2*B$  pages (on average)
  - use 1 page as an input buffer, 1 page as an output buffer
  - the remaining buffer pages are collectively referred to as the *current set*
- example - sort file in ascending order on some key  $k$ :
  - repeatedly pick record  $r$  in the current set and append it to the output buffer
  - keep output buffer sorted:  $r$ 's  $k$  value  $\geq$  largest  $k$  value in the output buffer
  - multiple such records in current set - choose the smallest one





- \* minimize the number of runs - optional
- example - sort file in ascending order on some key  $k$ :
  - use the extra space in the current set to bring in the next tuple from the input buffer
  - process all tuples in the input buffer, then read in the next page of the file
  - when the output buffer fills up, write it to disk (add its content to the run that is currently being built)
  - the current run is completed when every  $k$  value in the current set is  $<$  the largest  $k$  value in the output buffer; when this happens, the output buffer is written out (its content becomes the last page in the current run), and a new run is started

## Projection Based on Sorting

- improvement

- \* example

- pass 0:

- scan Exams: 1000 I/Os

- write out 250 pages:

- 20 available buffer pages

- 250 pages => 7 sorted runs about 40 pages long (except the last one, which is about 10 pages long)

- pass 1:

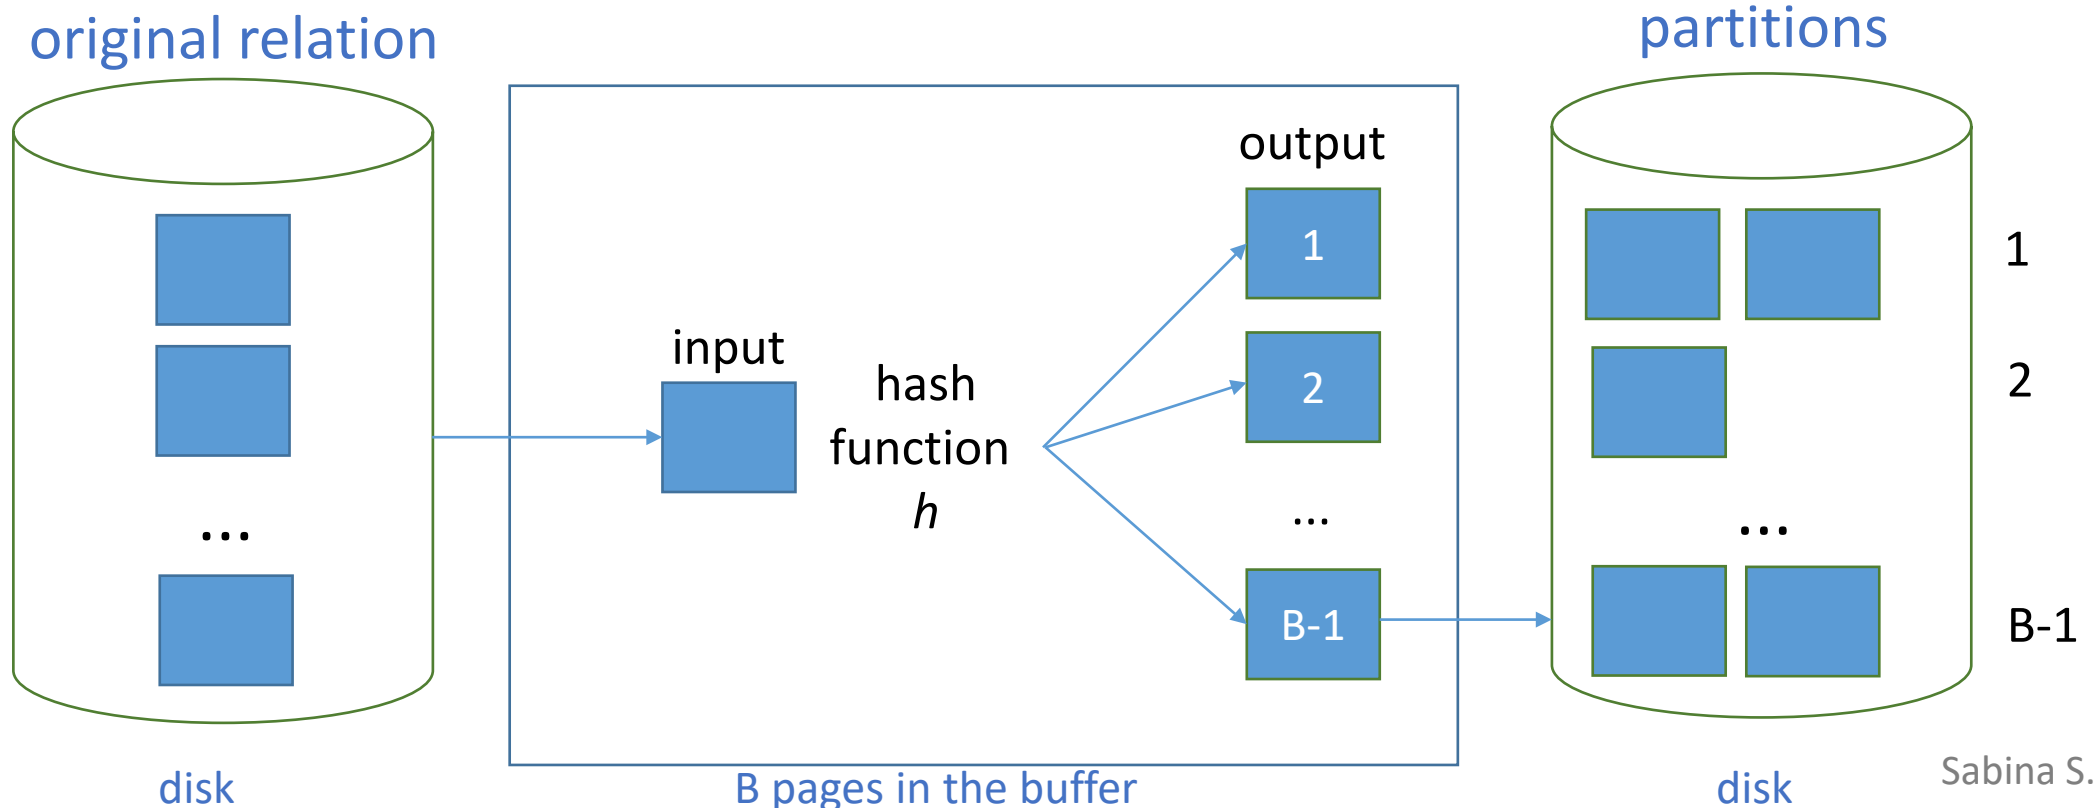
- read in all runs – cost: 250 I/Os

- merge runs

- total cost :  $1000 + 250 + 250 = 1500$  I/Os

## Projection Based on Hashing

- phases: partitioning & duplicate elimination
- partitioning phase:
  - 1 input buffer page – read in the relation one page at a time
  - hash function  $h$  – distribute tuples uniformly to one of  $B-1$  partitions
  - $B-1$  output buffer pages – one output page / partition



## Projection Based on Hashing

- partitioning phase:
  - read the relation using the input buffer page
  - for each tuple  $t$ :
    - discard unwanted fields  $\Rightarrow$  tuple  $t'$
    - apply hash function  $h$  to  $t'$
    - write  $t'$  to the output buffer page that it is hashed to by  $h$

$\Rightarrow$  B-1 partitions

- partition:
  - collection of tuples with:
    - common hash value
    - no unwanted fields
- tuples in different partitions are guaranteed to be distinct

## Projection Based on Hashing

- duplicate elimination phase:
  - process all partitions:
    - read in partition P, one page at a time
      - build in-memory hash table with hash function  $h_2 (\neq h)$  on all fields:
        - if a new tuple hashes to the same value as an existing tuple, compare them to check if they are distinct
        - eliminate duplicates
      - write duplicate-free hash table to result file
      - clear in-memory hash table
  - partition overflow
    - apply hash-based projection technique recursively (subpartitions)

## Projection Based on Hashing

- cost
  - partitioning:
    - read E: M I/Os
    - write E': T I/Os
  - duplicate elimination:
    - read in partitions: T I/Os

=> total cost:  $M + 2 * T$  I/Os
- Exams:
  - $1000 + 2 * 250 = 1500$  I/Os

## Set Operations

- intersection, cross-product
  - special cases of join (join condition for intersection - equality on all fields, no join condition for cross-product)
- union, set-difference
  - similar
- union:  $R \cup S$ 
  - sorting
    - sort R and S on all attributes
    - scan the sorted relations in parallel; merge them, eliminating duplicates
    - refinement
      - produce sorted runs of R and S, merge runs in parallel

## Set Operations

- union:  $R \cup S$ 
  - hashing
    - partition  $R$  and  $S$  with the same hash function  $h$
    - for each  $S$ -partition
      - build in-memory hash table (using  $h_2$ ) for the  $S$ -partition
      - scan corresponding  $R$ -partition, add tuples to hash table, discard duplicates
      - write out hash table
      - clear hash table



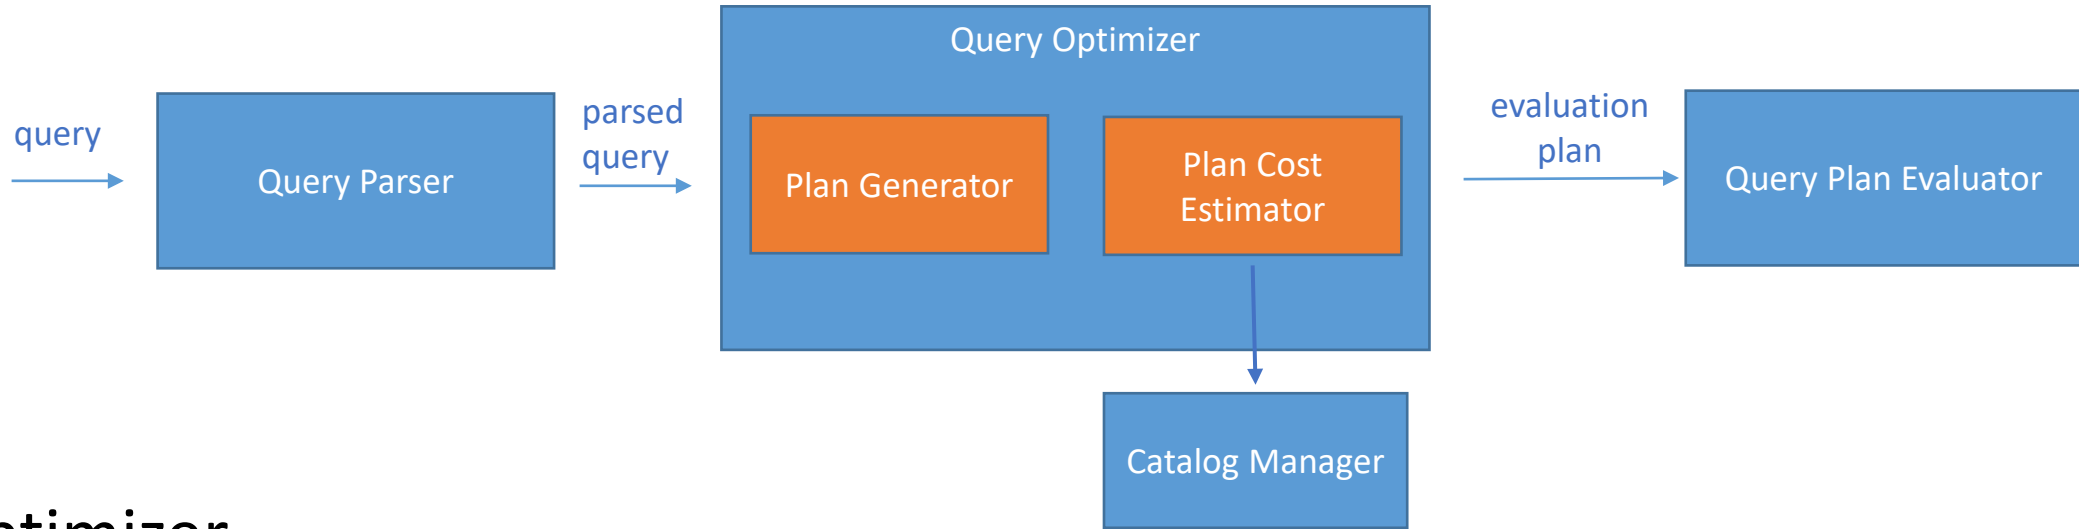
## Aggregate Operations

- without grouping
  - scan relation
  - maintain *running information* about scanned tuples
    - COUNT - count of values retrieved
    - SUM - *total* of values retrieved
    - AVG - *<total, count>* of values retrieved
    - MIN, MAX - smallest / largest value retrieved
- with grouping
  - sort relation on the grouping attributes
  - scan relation to compute aggregate operations for each group
  - improvement: combine sorting with aggregation computation
  - alternative approach based on hashing

## Aggregate Operations

- using existing indexes
  - index with a search key that includes all the attributes required by the query
    - work with the data entries in the index (instead of the data records)
  - attribute list in the GROUP BY clause is a prefix of the index search key (tree index)
    - get data entries (and records, if necessary) in the required order (i.e., avoid sorting)

# Query Optimization



- optimizer
  - objective
    - given a query  $Q$ , find a good evaluation plan for  $Q$
  - generates alternative plans for  $Q$ , estimates their costs, and chooses the one with the least estimated cost
  - uses information from the system catalogs


# Query Evaluation Plans

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      AND S.Age > 23
```

query

$\pi_{SName}(\sigma_{CID=7 \wedge Age>23}(Exams \otimes_{SID=SID} Students))$

relational algebra expression

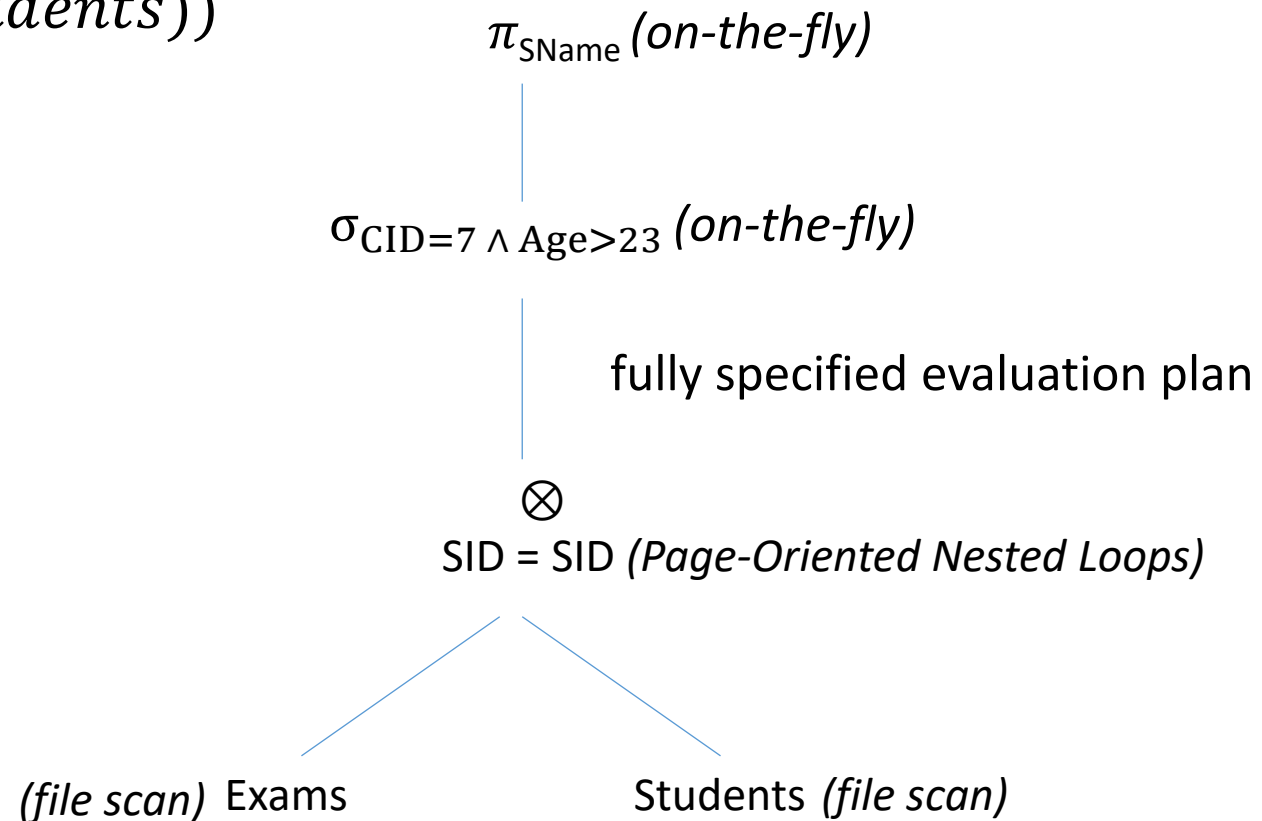


# Query Evaluation Plans

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      AND S.Age > 23
```

$\pi_{SName}(\sigma_{CID=7 \wedge Age>23}(Exams \otimes_{SID=SID} Students))$

- query evaluation plan
  - extended relational algebra tree
  - node – annotations
    - relation
      - access method
    - relational operator
      - implementation method

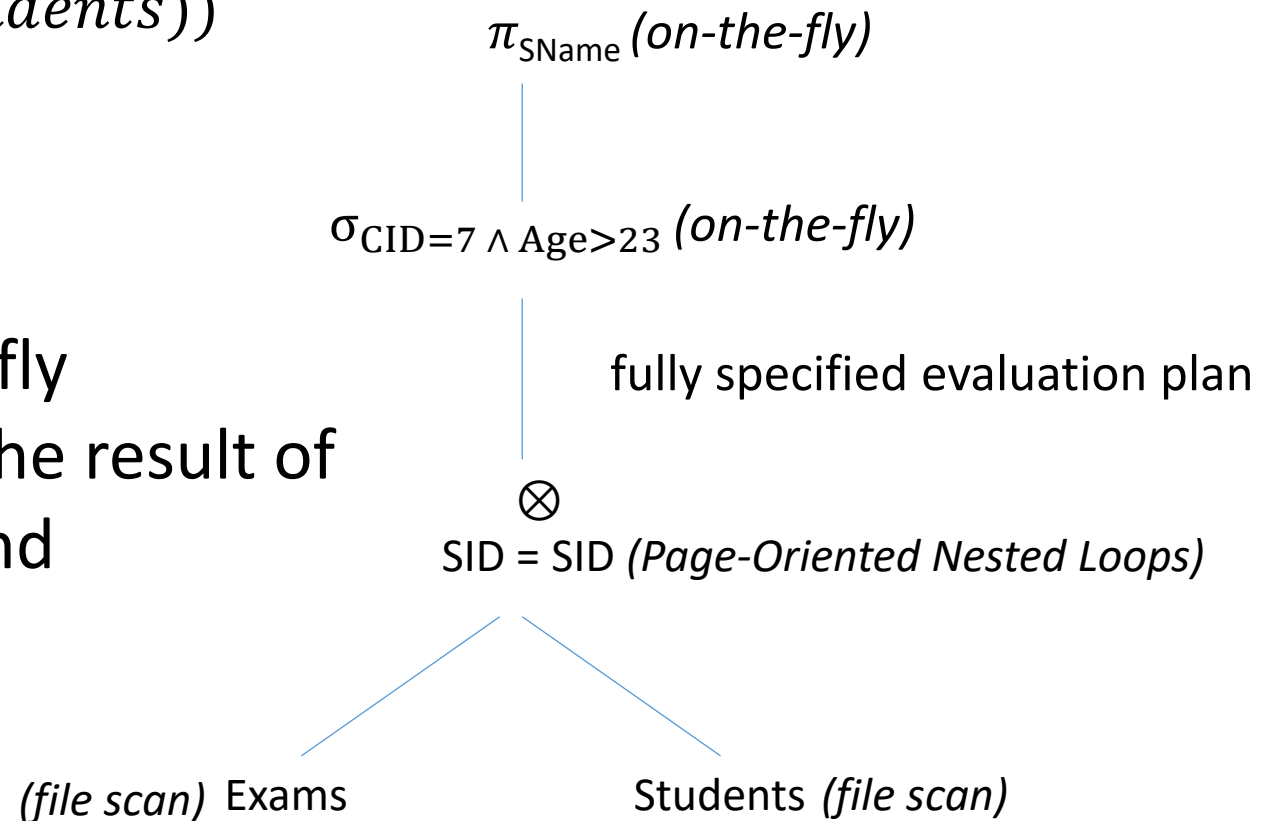


# Query Evaluation Plans

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      AND S.Age > 23
```

$\pi_{SName}(\sigma_{CID=7 \wedge Age>23}(Exams \otimes_{SID=SID} Students))$

- Page-Oriented Nested Loops Join
  - Exams – outer relation
- selection, projection applied on-the-fly to each tuple in the join result, i.e., the result of the join (before applying selection and projection) is not stored



# Pipelined Evaluation

SELECT \*

FROM Exams

WHERE  $\frac{EDate > '1-1-2020'}{T1}$  AND  $\frac{Grade > 8}{T2}$

$$\sigma_{Grade>8}(\sigma_{EDate>'1-1-2020'}(Exams))$$

- index / matches  $T1$
- $v1$  - *materialization*
  - evaluate  $T1$
  - write out result tuples to temporary relation  $R$ , i.e., tuples are *materialized*
  - apply the 2<sup>nd</sup> selection to  $R$
  - cost: read and write  $R$

# Pipelined Evaluation

SELECT \*

FROM Exams

WHERE  $\frac{\text{EDate} > '1-1-2020'}{T1}$  AND  $\frac{\text{Grade} > 8}{T2}$

- v2 – *pipelined evaluation*
  - apply the 2<sup>nd</sup> selection to each tuple in the result of the 1<sup>st</sup> selection as it is produced
  - i.e., 2<sup>nd</sup> selection operator is applied *on-the-fly*
  - saves the cost of writing out / reading in the temporary relation *R*



## Query Blocks – Units of Optimization

- parse  $Q \Rightarrow$  collection of query *blocks*  $\rightarrow$  passed on to the optimizer
- optimizer:
  - optimize one block at a time
- query *block* - SQL query:
  - without nesting
  - with exactly: one SELECT clause, one FROM clause
  - with at most: one WHERE clause, one GROUP BY clause, one HAVING clause
    - WHERE condition - CNF

# Query Blocks – Units of Optimization

- query Q:

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = (SELECT MAX(S2.Age)
               FROM Students S2)
GROUP BY S.SID
HAVING COUNT(*) > 2
```

nested block

- decompose query into a collection of blocks without nesting

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = Reference to nested block
GROUP BY S.SID
HAVING COUNT(*) > 2
```

# Query Blocks – Units of Optimization

## \* block optimization

- express query block as a relational algebra expression

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = Reference to nested block
GROUP BY S.SID
HAVING COUNT(*) > 2
```

$$\pi_{S.SID, MIN(E.EDate)}(\text{HAVING}_{COUNT(*) > 2}(\text{GROUP BY}_{S.SID}(\sigma_{S.SID = E.SID \wedge E.CID = C.CID \wedge C.Description = 'Elective' \wedge S.Age = value\_from\_nested\_block}(Students \times Exams \times Courses))))$$

- GROUP BY, HAVING – operators in the extended algebra used for plans
- argument list of projection can include aggregate operations

## Query Blocks – Units of Optimization

- query Q treated as a  $\sigma \pi \times$  algebra expression
- the remaining operations in Q are performed on the result of the  $\sigma \pi \times$  expression

```
SELECT S.SID, MIN(E.EDate)
FROM Students S, Exams E, Courses C
WHERE S.SID = E.SID AND E.CID = C.CID AND C.Description = 'Elective' AND
      S.Age = Reference to nested block
GROUP BY S.SID
HAVING COUNT(*) > 2
```

$$\pi_{S.SID, E.EDate}(\sigma_{S.SID = E.SID \wedge E.CID = C.CID \wedge C.Description = 'Elective' \wedge S.Age = value\_from\_nested\_block}(Students \times Exams \times Courses))$$

- attributes in GROUP BY, HAVING are added to the argument list of projection
- aggregate expressions in the argument list of projection are replaced by their argument attributes

## Query Blocks – Units of Optimization

### \* block optimization

- find best plan P for the  $\sigma \pi \times$  expression
- evaluate P  $\Rightarrow$  result set RS
- sort/hash RS  $\Rightarrow$  groups
- apply HAVING to eliminate some groups
- compute aggregate expressions in SELECT for each remaining group

$\pi_{S.SID, MIN(E.EDate)}($   
 $HAVING_{COUNT(*) > 2}($   
 $GROUP BY_{S.SID}($   
 $\pi_{S.SID, E.EDate}($   
 $\sigma_{S.SID = E.SID \wedge E.CID = C.CID \wedge C.Description = 'Elective' \wedge S.Age = value\_from\_nested\_block}($   
 $Students \times Exams \times Courses ))))$

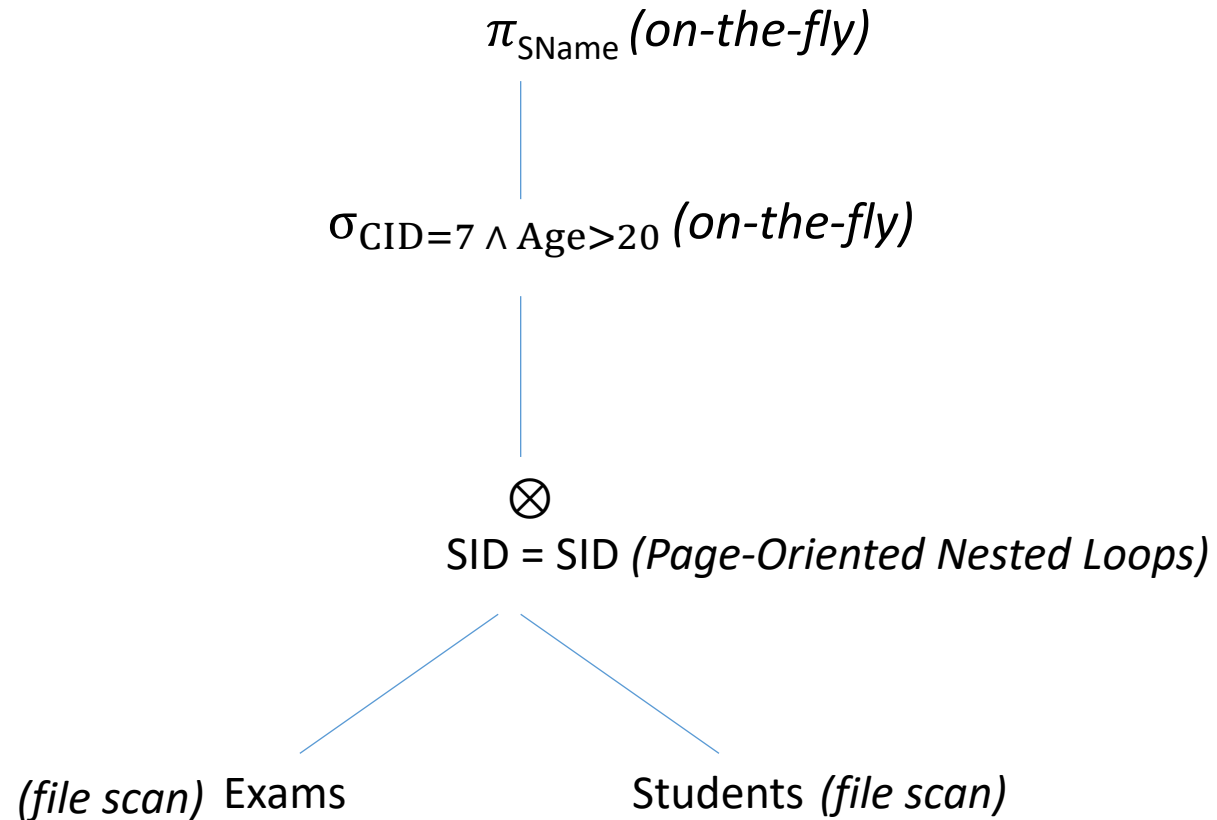
## Motivating Example

\* E - 1000 pages \*

\* S - 500 pages \*

```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      AND S.Age > 20
```

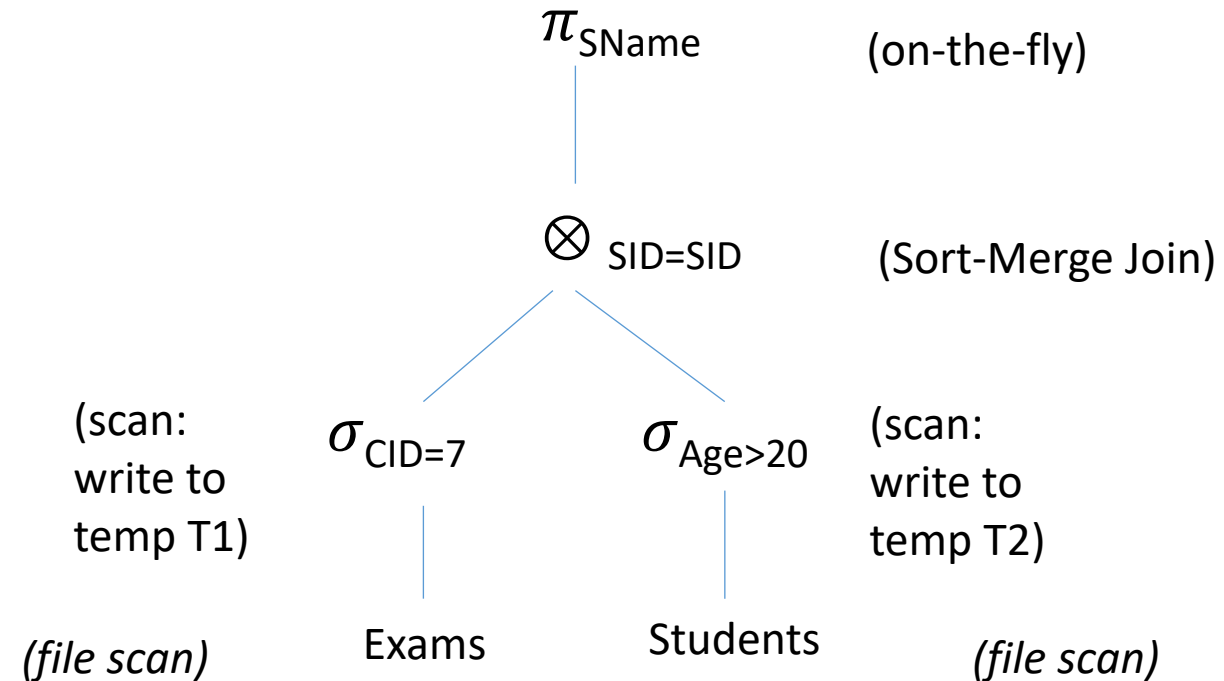
- $\sigma, \pi$  – on-the-fly
- cost of plan – very high:
  - $1000 + 1000 * 500 = 501,000$  I/Os



# Motivating Example

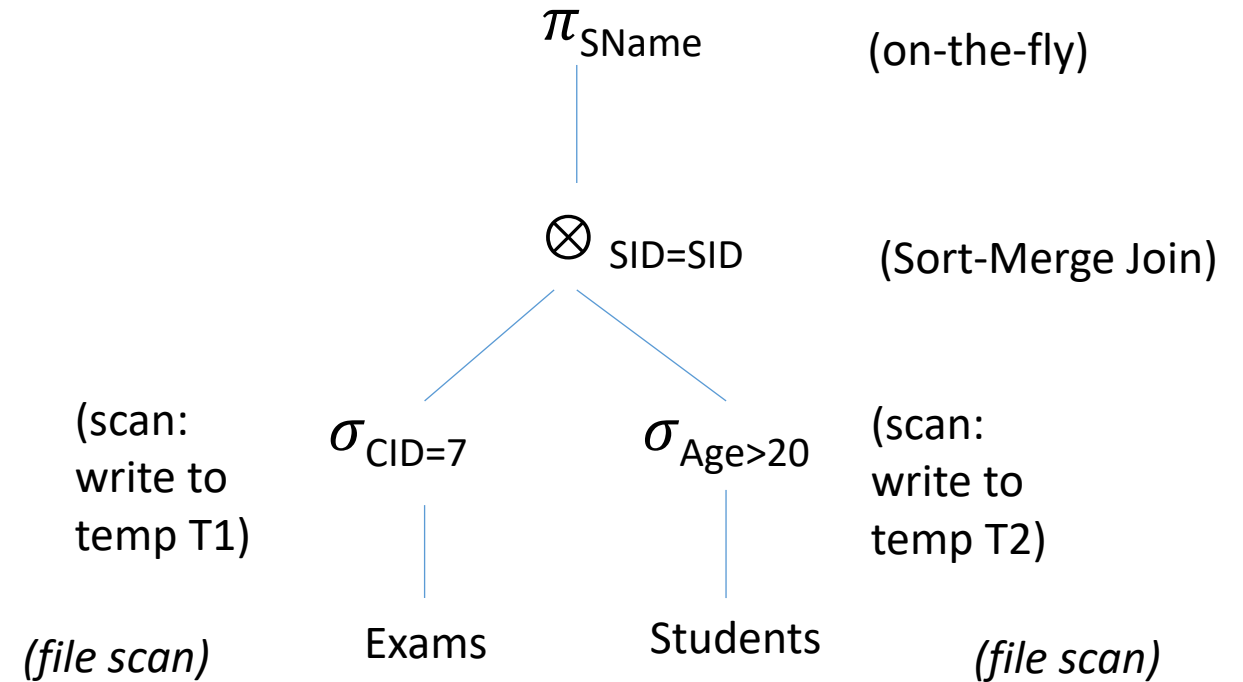
## \* optimizations

- reduce sizes of the relations to be joined
  - push selections, projections ahead of the join
- alternative plans
  - push selections ahead of joins
- selection
  - file scan
  - write the result to a temporary relation on disk
- join the temporary relations using Sort-Merge Join



## Motivating Example

- 5 available buffer pages
- cost
  - $\sigma_{CID=7}$ 
    - scan Exams: 1000 I/Os
    - write T1
      - assume exams are uniformly distributed across all courses, i.e., T1 has 10 pages (there are 100 courses)
  - $\sigma_{Age>20}$ 
    - scan Students: 500 I/Os
    - write T2
      - assume ages are uniformly distributed over the range 19 to 22, i.e., T2 has 250 pages

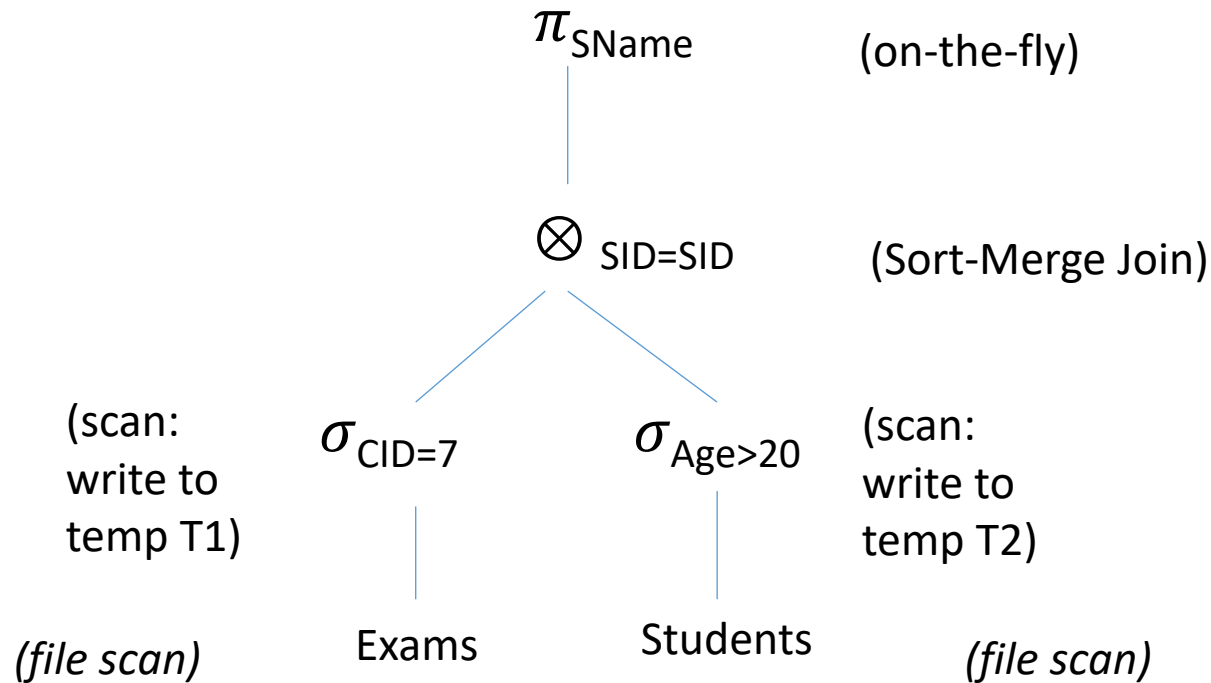




## Motivating Example

- 5 available buffer pages
- cost
  - Sort-Merge Join
    - T1 - 10 pages
      - sort T1:  $2 * 2 * 10 = 40$  I/Os
    - T2 - 250 pages
      - sort T2:  $2 * 4 * 250 = 2000$  I/Os
    - merge sorted T1 and T2
      - $10 + 250 = 260$  I/Os
  - $\pi$  - on the fly

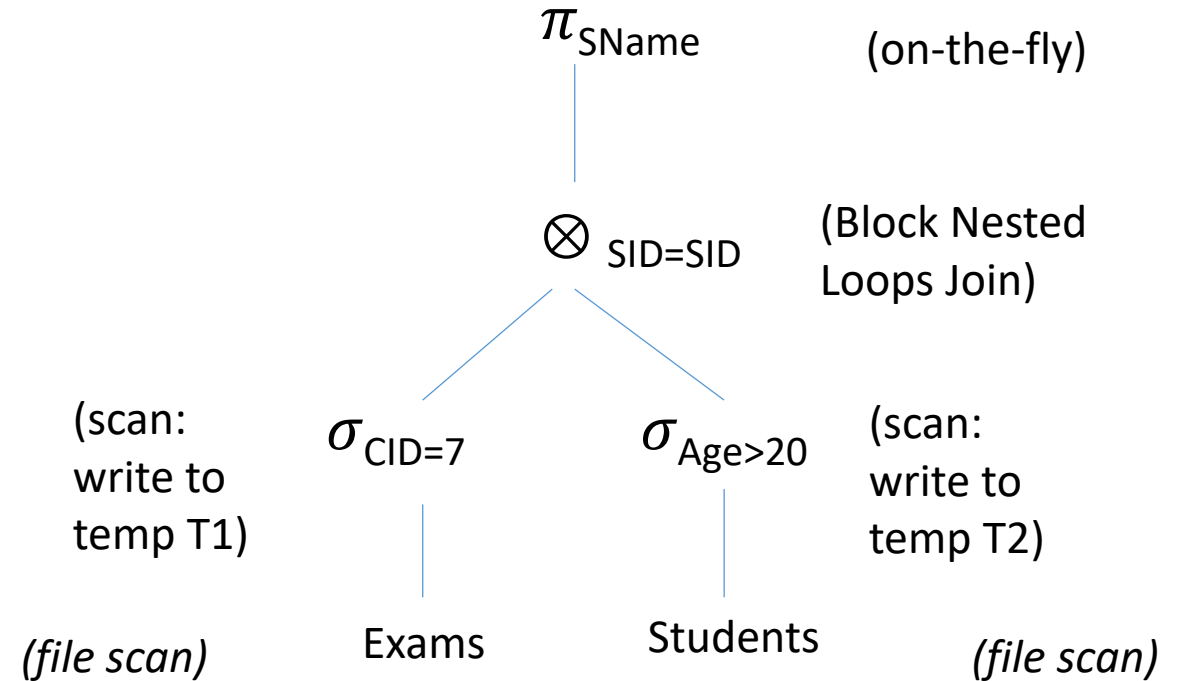
=> **total cost:**  $\underbrace{1000 + 10 + 500 + 250}_{\text{selection}} + \underbrace{40 + 2000 + 260}_{\text{join}} = 4060 \text{ I/Os}$



## Motivating Example

- 5 available buffer pages
- cost
  - Block Nested Loops Join
    - T1 - 10 pages, T2 - 250 pages
    - T1 - outer relation
      - => scan T1: 10 I/Os
    - $\lceil 10/3 \rceil = 4$  T1 blocks
      - => T2 scanned 4 times:  $4 * 250 = 1000$  I/Os
    - BNLJ cost:  $10 + 1000 = 1010$  I/Os
  - $\pi$  - on the fly

=> **total cost:**  $\underbrace{1000 + 10 + 500 + 250}_{\text{selection}} + \underbrace{10 + 1000}_{\text{join}} = \mathbf{2770 \text{ I/Os}}$



## Motivating Example

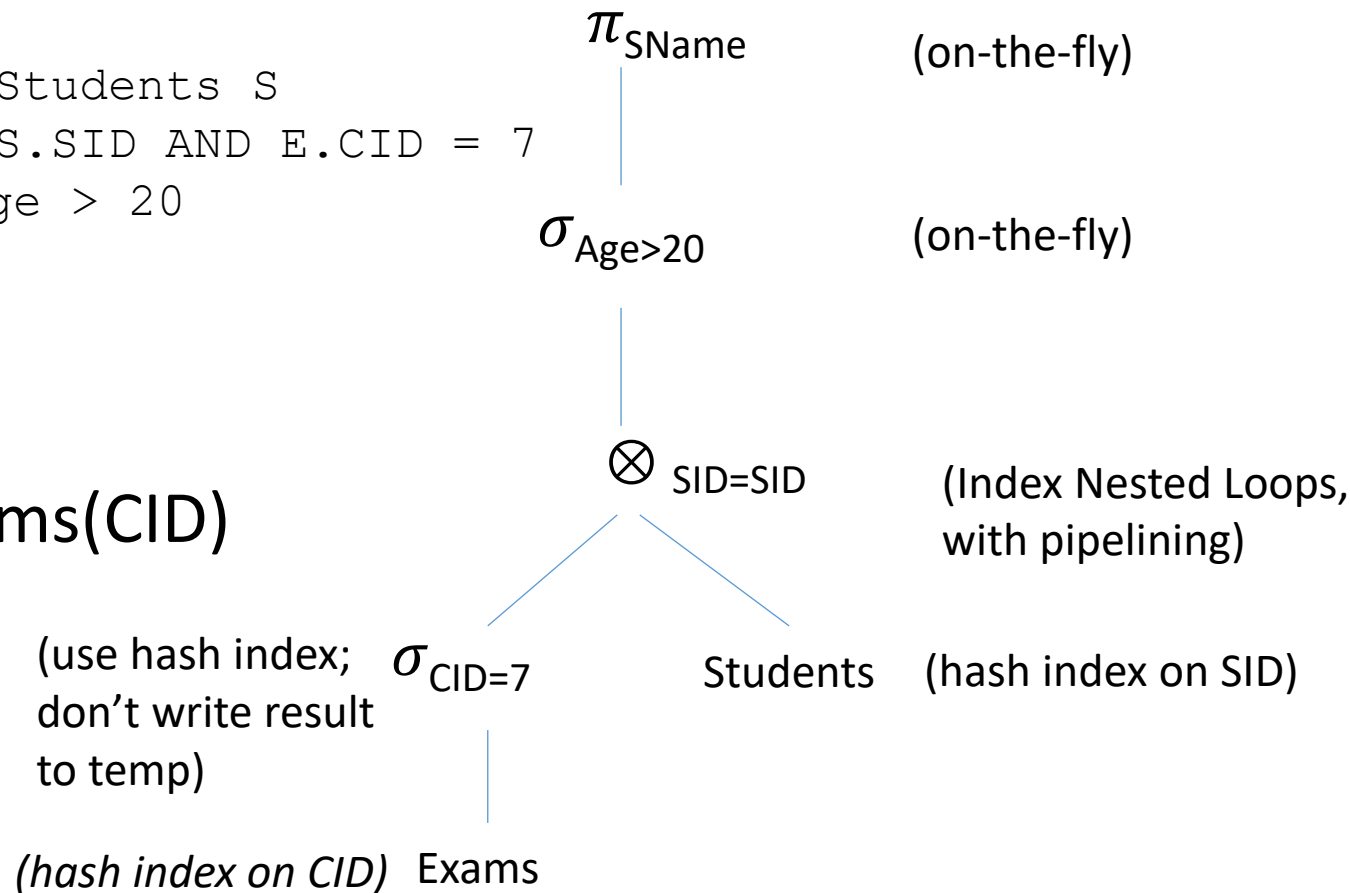
- push projections ahead of joins
  - drop unwanted columns while scanning Exams and Students to evaluate selections => T1[SID], T2[SID, SName]
- T1 fits within 3 buffer pages
  - => T2 scanned only once
  - => **total cost**: about **2000 I/Os**

# Motivating Example

## \* optimizations

- investigate the use of indexes
- clustered static hash index on Exams(CID)
- hash index on Students(SID)
- cost
  - $\sigma_{CID=7}$ 
    - assume exams are uniformly distributed across all courses => 100,000 exams / 100 courses => 1,000 exams / course
    - clustered index on CID => 1,000 tuples for course with CID=7 appear consecutively within the same bucket => cost: 10 I/Os
    - the result of the selection is not materialized, the join is pipelined

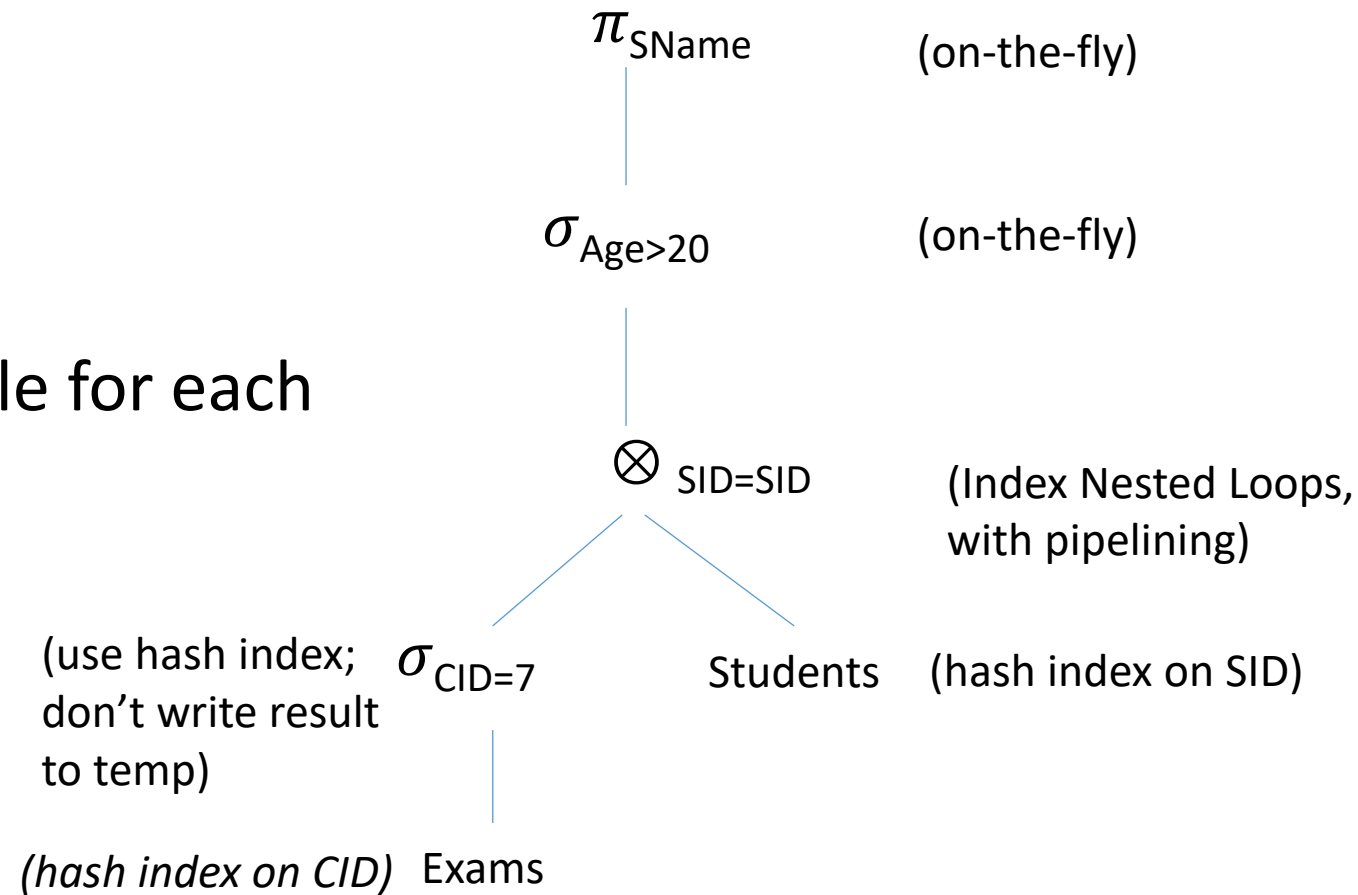
```
SELECT S.SName
FROM Exams E, Students S
WHERE E.SID = S.SID AND E.CID = 7
      AND S.Age > 20
```



## Motivating Example

- cost
  - Index Nested Loops
    - find matching Students tuple for each selected exam
  - use hash index on SID
    - assume the index uses a1 => cost of 1.2 I/Os (on avg.) per exam
- $\sigma, \pi$  – performed on-the-fly on each tuple in the result of the join

$$\Rightarrow \text{total cost} = \underbrace{10}_{\sigma \text{ on Exams}} + \underbrace{1000}_{\text{num. of Exams tuples}} * \underbrace{1.2}_{\text{find matching Students tuple (on avg.)}} = \mathbf{1210 \text{ I/Os}}$$



\* can we push the selection *Age>20* ahead of the join?

# References

- [Ra02] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems (3rd Edition), McGraw-Hill, 2002
- [Da03] DATE, C.J., An Introduction to Database Systems (8<sup>th</sup> Edition), Addison-Wesley, 2003
- [Ga09] GARCIA-MOLINA, H., ULLMAN, J., WIDOM, J., Database Systems: The Complete Book (2nd Edition), Pearson Education, 2009
- [Ra02S] RAMAKRISHNAN, R., GEHRKE, J., Database Management Systems, Slides for the 3<sup>rd</sup> Edition,  
<http://pages.cs.wisc.edu/~dbbook/openAccess/thirdEdition/slides/slides3ed.html>
- [Si19] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts (7th Edition), McGraw-Hill, 2019
- [Si19S] SILBERSCHATZ, A., KORTH, H., SUDARSHAN, S., Database System Concepts, Slides for the 7th Edition, <http://codex.cs.yale.edu/avi/db-book/>
- [Ul11] ULLMAN, J., WIDOM, J., A First Course in Database Systems,  
<http://infolab.stanford.edu/~ullman/fcdb.html>